

# Lab 11-1 Convolution Neural Network & Data Pipelines

NTHU DataLab, 2025

# Outline

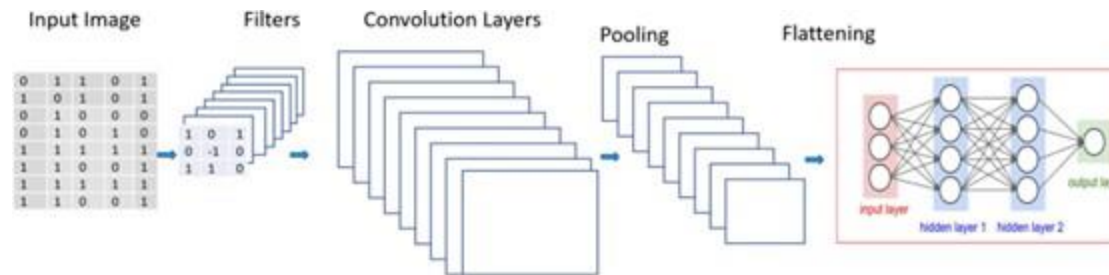
- Convolution neural network
- Input pipeline
- Optimization for Input pipeline

# Outline

- Convolution neural network
- Input pipeline
- Optimization for Input pipeline

# Convolution Neural Network

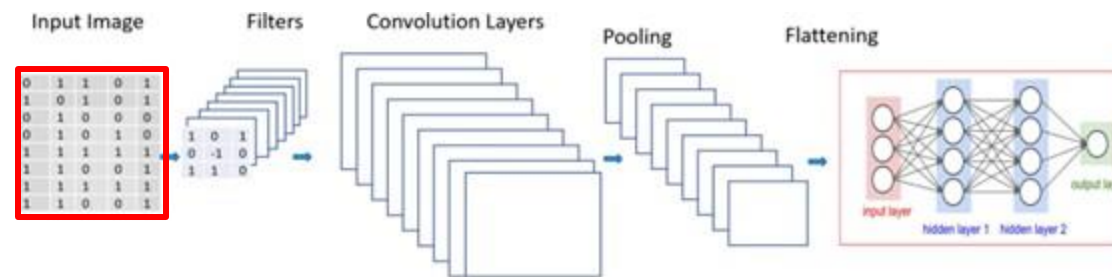
- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers



```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```

# Convolution Neural Network

- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers

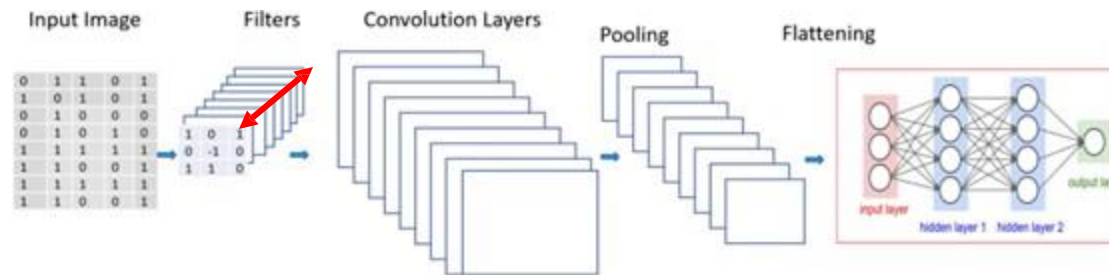


Width \* Height \* Channel

```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```

# Convolution Neural Network

- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers

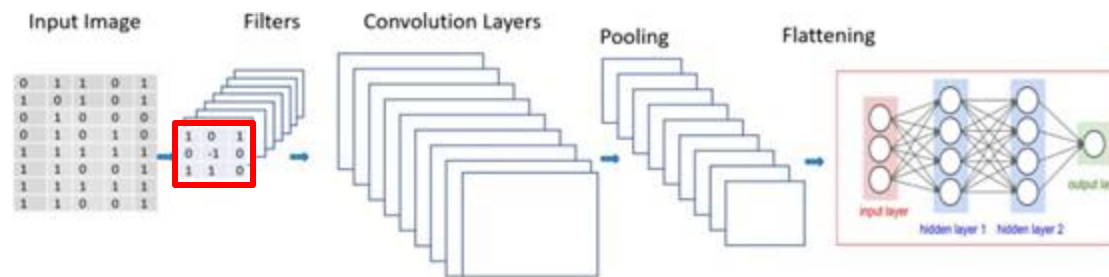


Filters: the number of output filters in the convolution

```
model = models.Sequential()
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))
model.add(layers.MaxPooling2D((2, 2)))
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```

# Convolution Neural Network

- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers



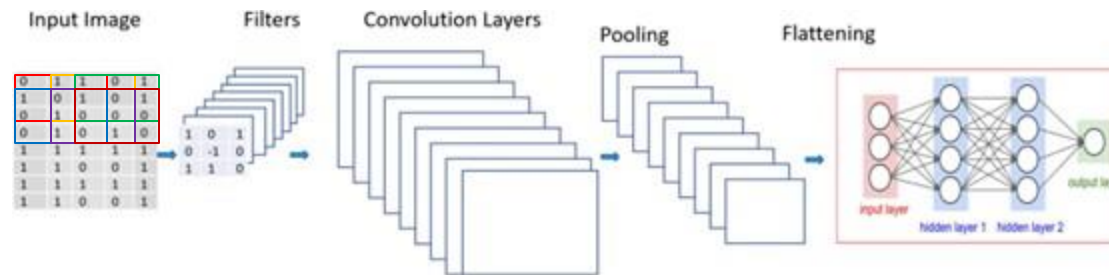
Kernel size: specifying the height and width of the 2D convolution window

```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```



# Convolution Neural Network

- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers



Layer (type)	Output Shape
conv2d_28 (Conv2D)	(None, 26, 26, 32)
max_pooling2d_12 (MaxPooling2D)	(None, 13, 13, 32)
conv2d_29 (Conv2D)	(None, 6, 6, 64)
max_pooling2d_13 (MaxPooling2D)	(None, 3, 3, 64)
conv2d_30 (Conv2D)	(None, 3, 3, 64)

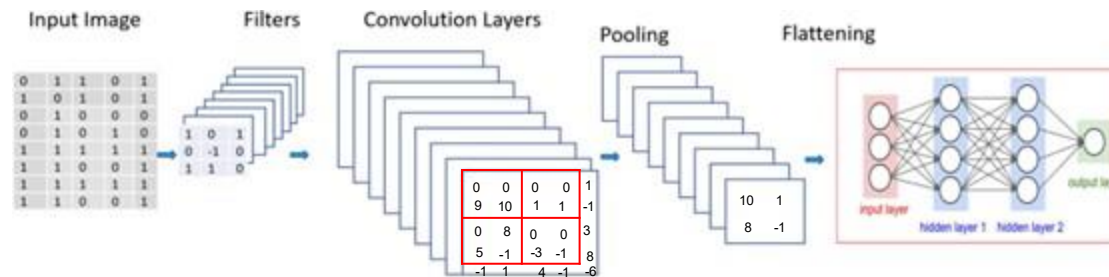
```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```

(Default value)



# Convolution Neural Network

- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers

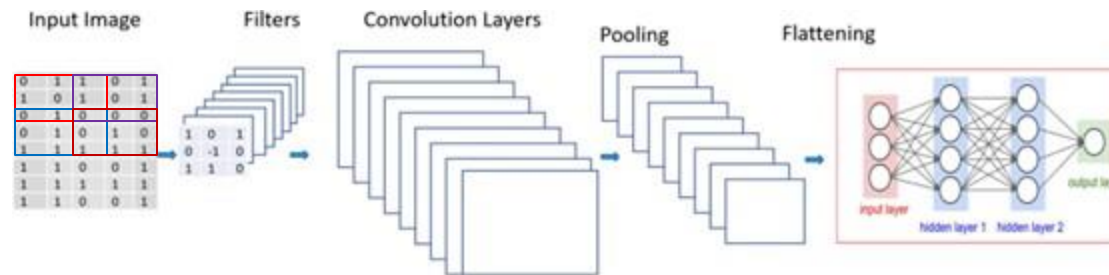


Layer (type)	Output Shape
conv2d_28 (Conv2D)	(None, 26, 26, 32)
max_pooling2d_12 (MaxPooling2D)	(None, 13, 13, 32)
conv2d_29 (Conv2D)	(None, 6, 6, 64)
max_pooling2d_13 (MaxPooling2D)	(None, 3, 3, 64)
conv2d_30 (Conv2D)	(None, 3, 3, 64)

```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```

# Convolution Neural Network

- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers

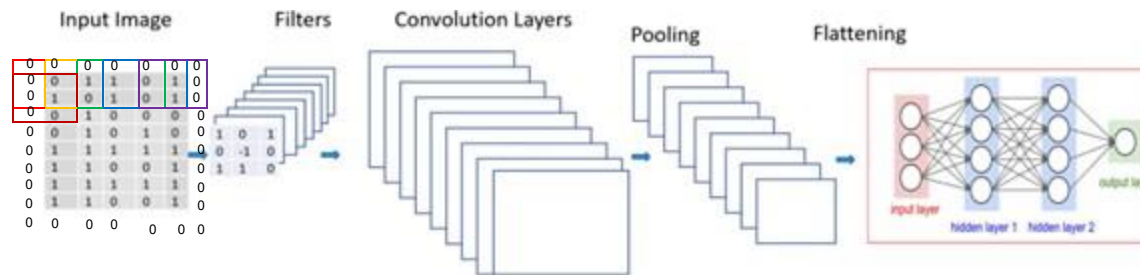


Layer (type)	Output Shape
conv2d_28 (Conv2D)	(None, 26, 26, 32)
max_pooling2d_12 (MaxPooling2D)	(None, 13, 13, 32)
conv2d_29 (Conv2D)	(None, 6, 6, 64)
max_pooling2d_13 (MaxPooling2D)	(None, 3, 3, 64)
conv2d_30 (Conv2D)	(None, 3, 3, 64)

```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```

# Convolution Neural Network

- Build a CNN model via Sequential API
  - A stack of Conv2D and MaxPooling2D layers



Layer (type)	Output Shape
conv2d_28 (Conv2D)	(None, 26, 26, 32)
max_pooling2d_12 (MaxPooling2D)	(None, 13, 13, 32)
conv2d_29 (Conv2D)	(None, 6, 6, 64)
max_pooling2d_13 (MaxPooling2D)	(None, 3, 3, 64)
conv2d_30 (Conv2D)	(None, 3, 3, 64)

```
model = models.Sequential()  
model.add(layers.Conv2D(32, (3, 3), strides=(1,1), padding='valid', activation='relu', input_shape=(28, 28, 1)))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(2,2), padding='valid', activation='relu'))  
model.add(layers.MaxPooling2D((2, 2)))  
model.add(layers.Conv2D(64, (3, 3), strides=(1,1), padding='same', activation='relu'))
```

When padding="same" and strides=1, the output has the same size as the input.

# Outline

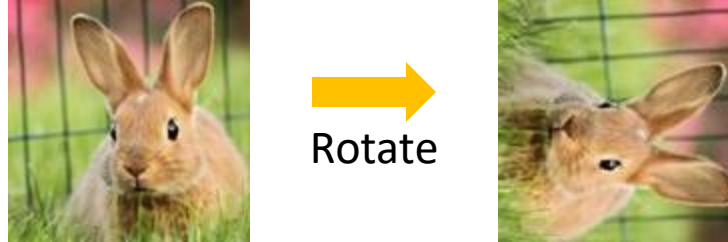
- Convolution neural network
- **Input pipeline**
- Optimization for Input pipeline

# Input Pipeline

- A series of input data processing before training

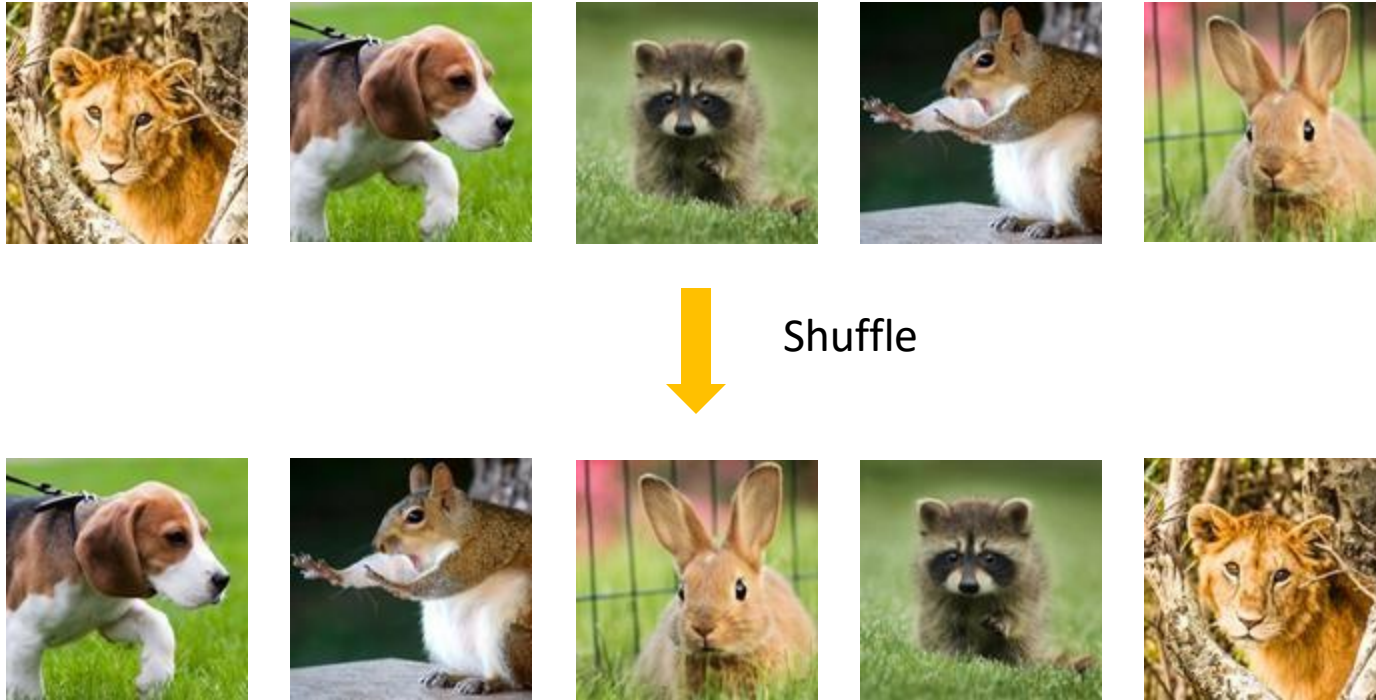
# Input Pipeline

- A series of input data processing before training



# Input Pipeline

- A series of input data processing before training



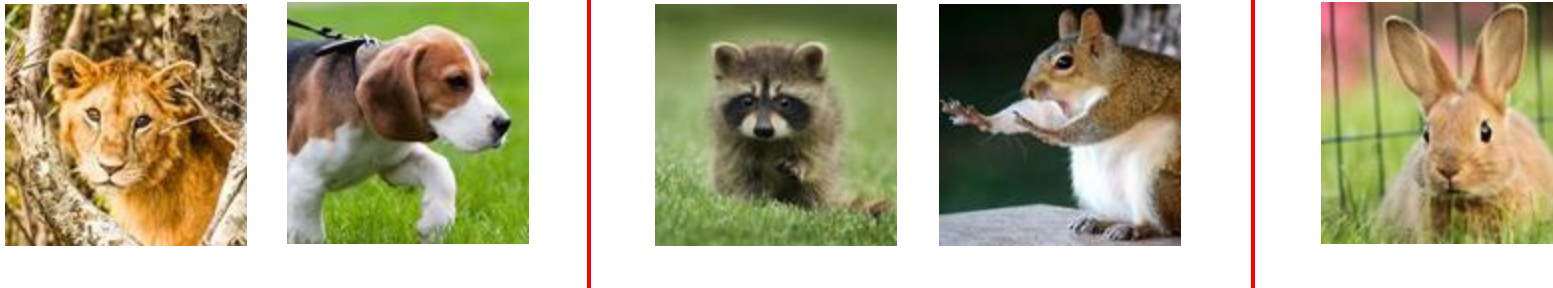


# Input Pipeline

- A series of input data processing before training



Batch



# Input Pipeline

- A series of input data processing before training
- Building the input pipeline is long and painful, and it's hard to reuse due to different type of data
- TensorFlow provides an API `tf.data` enables you to build complex input pipelines from simple, reusable pieces

# Input Pipeline - TensorFlow

- To apply transformations on your input data, we will need to construct a `tf.data.Dataset` object

# Input Pipeline - TensorFlow

- Construct Dataset
  - For small data (**in-memory**)

```
raw_data_a: [[0.59004802, 0.68869704, 0.67771658, 0.25277111, 0.44878355],  
             [0.2194635 , 0.23323033, 0.37668097, 0.0523581 , 0.84413446],  
             ⋮  
             [0.2194635 , 0.23323033, 0.37668097, 0.0523581 , 0.84413446],  
             [0.94882014, 0.3818479 , 0.93550471, 0.23102154, 0.66095901]]  
  
raw_data_b: [ 0, 1, 2, ..., 199]
```

```
# number of samples  
n_samples = 200  
  
# an array with shape (n_samples, 5)  
raw_data_a = np.random.rand(n_samples, 5)  
# a list with length of n_samples from 0 to n_samples-1  
raw_data_b = np.arange(n_samples)  
  
# this tells the dataset that each row of raw_data_a is corresponding to each element of raw_data_b  
raw_dataset = tf.data.Dataset.from_tensor_slices((raw_data_a, raw_data_b))  
print(raw_dataset)  
  
<TensorSliceDataset shapes: ((5,), ()), types: (tf.float64, tf.int64)>
```

All input tensors must have the same size in their first dimensions

The given tensors are sliced along their **first dimension**.

# Input Pipeline - TensorFlow

- Construct Dataset
  - For small data (**in-memory**)

```
# number of samples
n_samples = 200

# an array with shape (n_samples, 5)
raw_data_a = np.random.rand(n_samples, 5)
# a list with length of n_samples from 0 to n_samples-1
raw_data_b = np.arange(n_samples)

# this tells the dataset that each row of raw_data_a is corresponding to each element of raw_data_b
raw_dataset = tf.data.Dataset.from_tensor_slices((raw_data_a, raw_data_b))
print(raw_dataset)
```

```
<TensorSliceDataset shapes: ((5,), ()), types: (tf.float64, tf.int64)>
```

```
train_ds = tf.data.Dataset.from_tensor_slices(data label (x_train, y_train))
test_ds = tf.data.Dataset.from_tensor_slices((x_test, y_test))
```

# Input Pipeline - TensorFlow

- Transformations
  - **Map**: apply the function to each the elements of this dataset
  - **Shuffle**: maintains a fixed-size buffer and chooses the next element uniformly at random from that buffer
  - **Batch**: combines consecutive elements of this dataset into batches
  - **Repeat**: repeat this dataset so each original value is seen multiple times
  - **Prefetch**: allows later elements to be prepared while the current element is being processed

# Input Pipeline - TensorFlow

- Transformations
  - **Map**: apply the function to each the elements of this dataset

```
def preprocess_function(one_row_a, one_b):  
    """  
        Input: one slice of the dataset  
        Output: modified slice  
    """  
  
    # Do some data preprocessing, you can also input filenames and load data in here  
    # Here, we transform each row of raw_data_a to its sum and mean  
    one_row_a = [tf.reduce_sum(one_row_a), tf.reduce_mean(one_row_a)]  
  
    return one_row_a, one_b  
  
raw_dataset = raw_dataset.map(preprocess_function)  
print(raw_dataset)
```

Map function

```
<MapDataset shapes: ((2,), ()), types: (tf.float64, tf.int64)>
```

reduce the dimension



# Input Pipeline - TensorFlow

- Transformations
  - **Map:** apply the function to each the elements of this dataset
    - Data augmentation: a technique to increase the diversity of your training set by applying random transformations such as image rotation

```
def pre_train_data(img, label):  
    distorted_img = tf.image.random_crop(img, [IMAGE_SIZE_CROPPED, IMAGE_SIZE_CROPPED, IMAGE_DEPTH])  
    distorted_img = tf.image.random_flip_left_right(distorted_img)  
    distorted_img = tf.image.random_brightness(distorted_img, max_delta=63)  
    distorted_img = tf.image.random_contrast(distorted_img, lower=0.2, upper=1.8)  
    distorted_img = tf.image.per_image_standardization(distorted_img)  
  
    return distorted_img, label
```

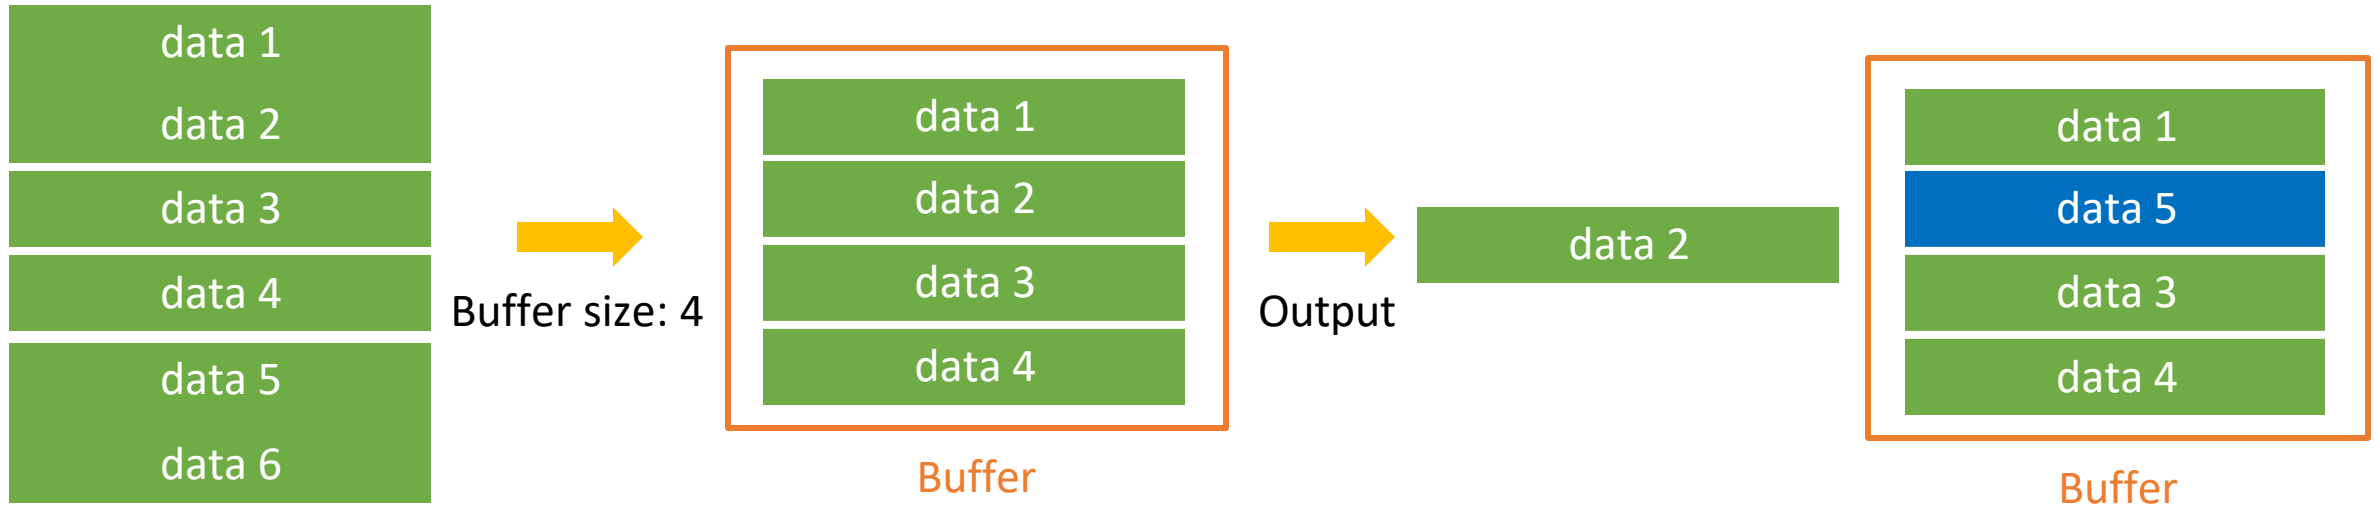


# Input Pipeline - TensorFlow

- Transformations
  - **Shuffle**: maintains a fixed-size buffer and chooses the next element uniformly at random **from that buffer**

Buffer size

```
dataset = raw_dataset.shuffle(16)
```



# Input Pipeline - TensorFlow

- Transformations
  - **Batch**: combines consecutive elements of this dataset into batches

```
dataset = dataset.batch(2, drop_remainder=False)
```

- Be careful that if you apply shuffle after batch, you'll get shuffled batch but data in a batch remains the same



shuffle -> batch

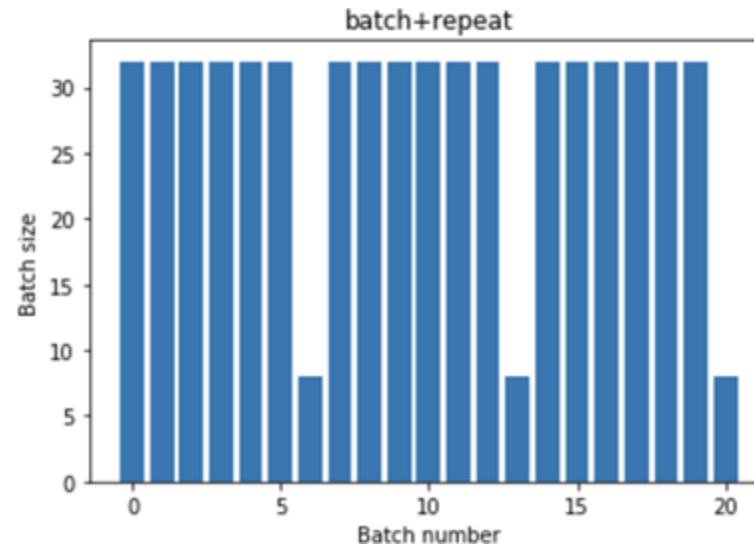
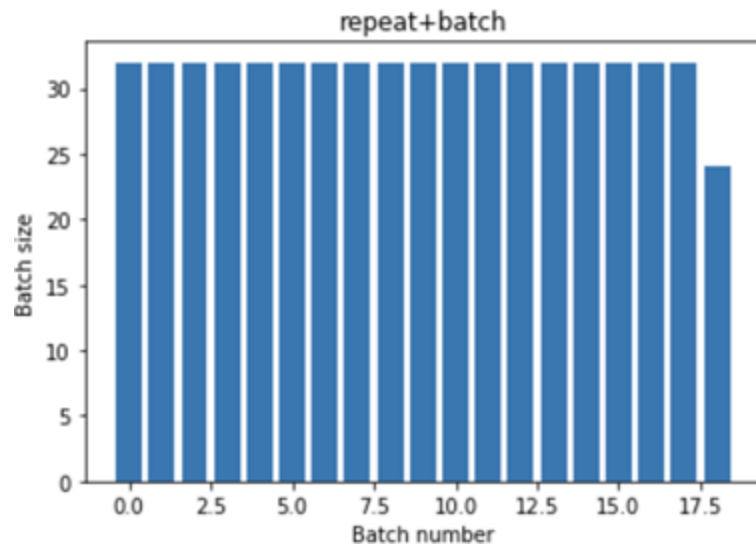


batch -> shuffle

# Input Pipeline - TensorFlow

- Transformations
  - **Repeat:** repeat this dataset so each original value is seen multiple times

```
dataset = dataset.repeat(2)
```



# Input Pipeline - TensorFlow

- Transformations
  - **Prefetch**: allows later elements to be prepared while the current element is being processed
  - This often improves latency and throughput, at the cost of using additional memory to store prefetched elements

```
dataset = dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
```

it will prompt the `tf.data` runtime to tune the value dynamically at runtime

# Input Pipeline - TensorFlow

- Transformations
  - **Prefetch**: allows later elements to be prepared while the current element is being processed
  - This often improves latency and throughput, at the cost of using additional memory to store prefetched elements

```
dataset = dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
```

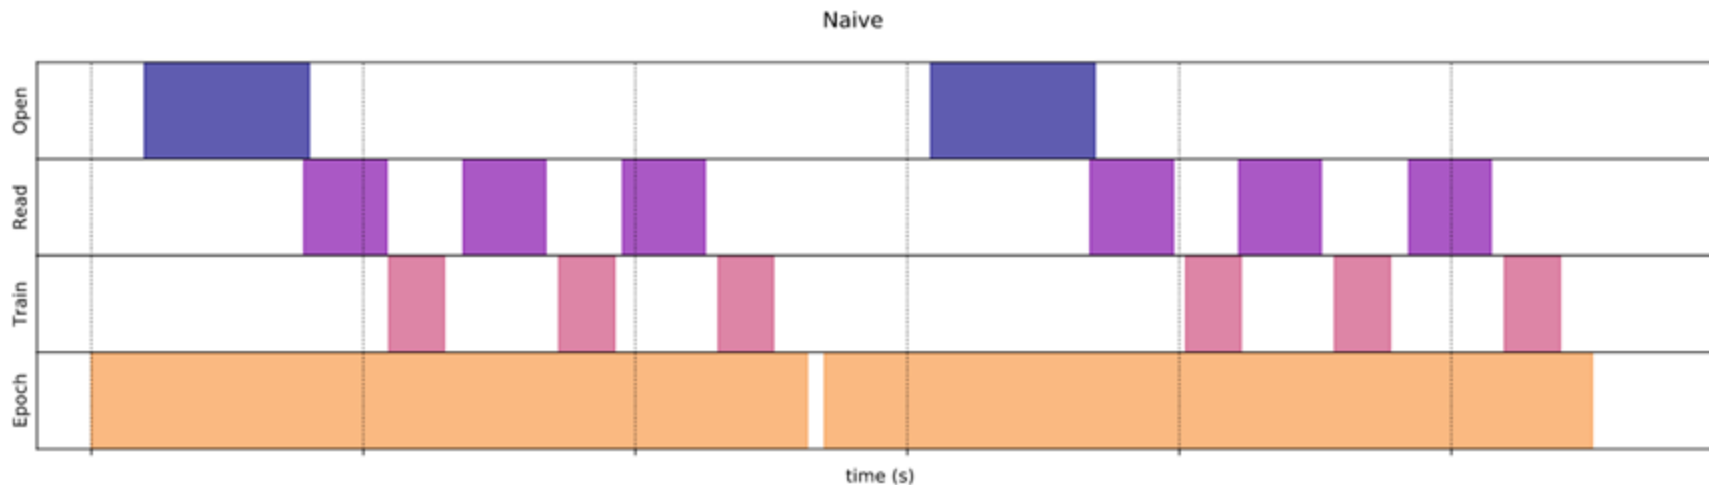
```
def preprocess_function(one_row_a, one_b):  
    """  
        Input: one slice of the dataset  
        Output: modified slice  
    """  
    # Do some data preprocessing, you can also input filenames and load data in here  
    # Here, we transform each row of raw_data_a to its sum and mean  
    one_row_a = [tf.reduce_sum(one_row_a), tf.reduce_mean(one_row_a)]  
  
    return one_row_a, one_b  
raw_dataset = raw_dataset.map(preprocess_function, num_parallel_calls=tf.data.experimental.AUTOTUNE)
```

# Input Pipeline - TensorFlow

- Transformations
  - **Prefetch**: allows later elements to be prepared while the current element is being processed

```
dataset = dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
```

it will prompt the `tf.data` runtime to tune the value dynamically at runtime



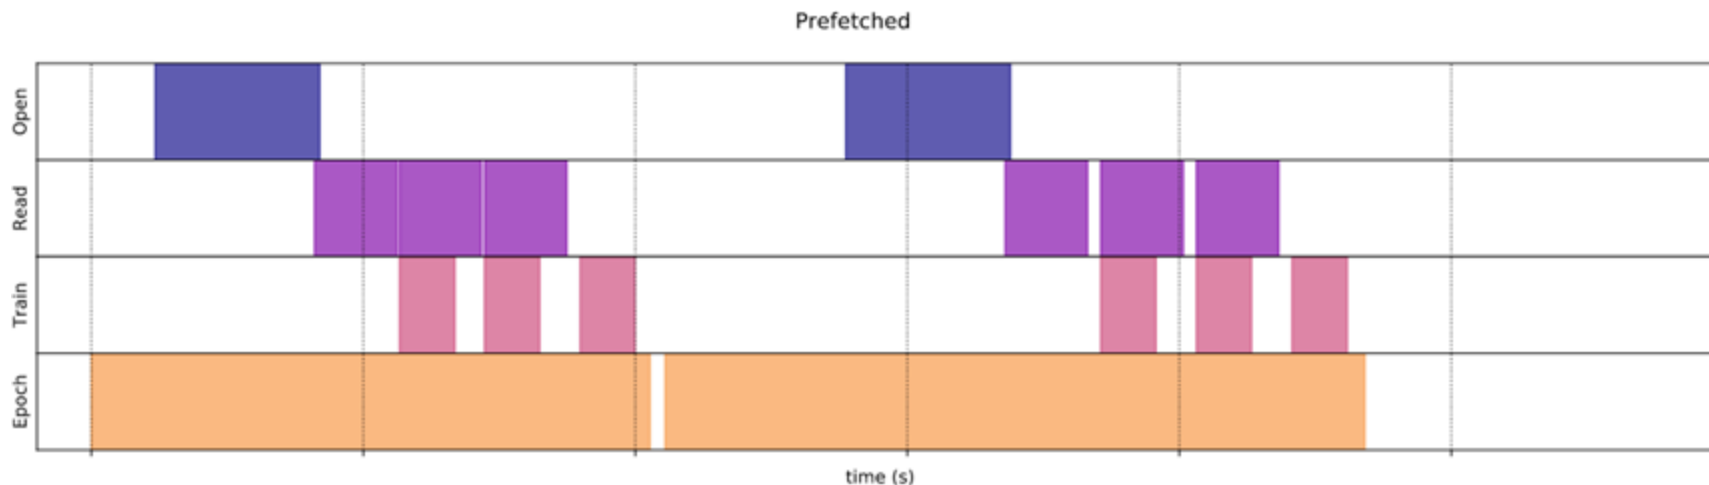


# Input Pipeline - TensorFlow

- Transformations
  - **Prefetch**: allows later elements to be prepared while the current element is being processed

```
dataset = dataset.prefetch(buffer_size=tf.data.experimental.AUTOTUNE)
```

it will prompt the `tf.data` runtime to tune the value dynamically at runtime



# Input Pipeline - TensorFlow

- Now you can iterate through the data and train
- Aware that if you do batch, the first dimension will be the batch size

```
for img, label in dataset_train.take(1):  
    print(img.shape)  
    print(label.shape)
```

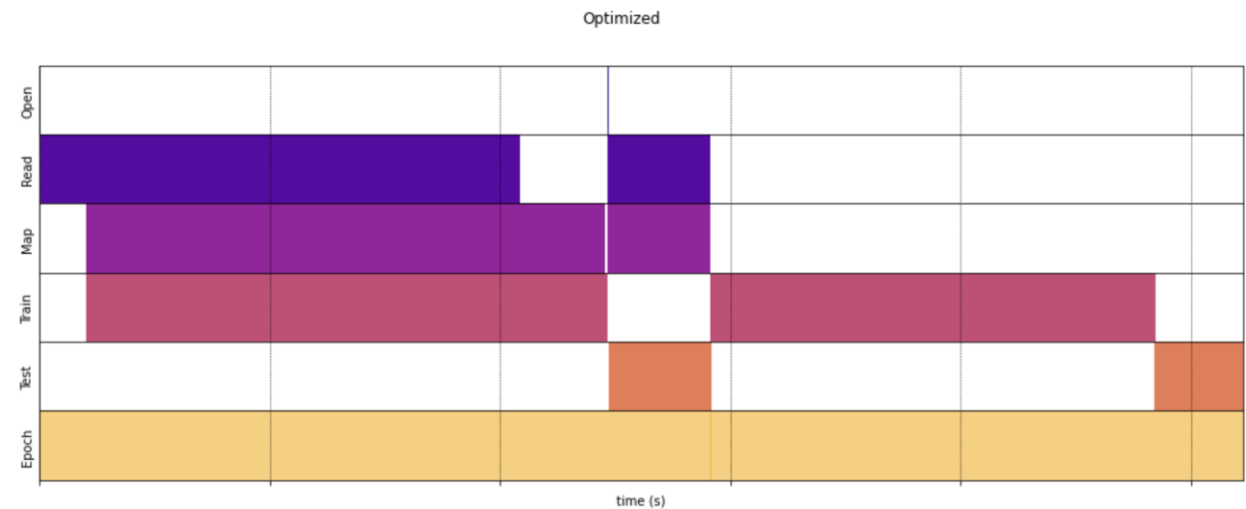
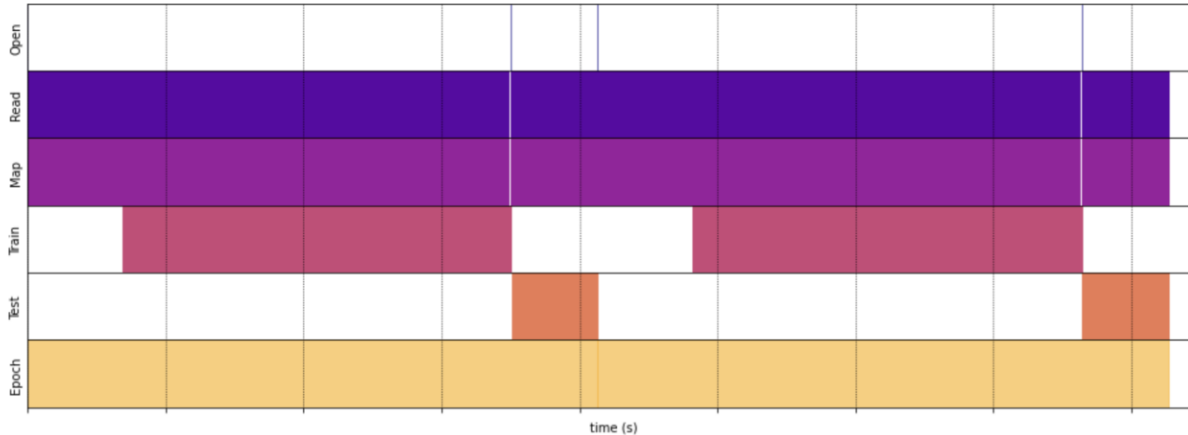
```
(32, 224, 224, 3)  
(32,)
```

# Outline

- Convolution neural network
- Input pipeline
- Optimization for Input pipeline

# Optimization for Input pipeline

1. **prefetching**: overlaps the preprocessing and model execution of a training step.
2. **Interleave (Parallelizing data extraction)**: parallelize the data loading step, interleaving the contents of other datasets (such as data file readers).
3. **Parallel mapping**: parallelized mapping across multiple CPU cores.
4. **Caching**: cache a dataset, save some operations (like file opening and data reading) from being executed during each epoch.
5. **Vectorizing mapping**: batch before map, so that mapping can be vectorized.



# Optimization for Input pipeline

```
dataset_train_optimized = tf.data.Dataset.range(1).interleave(dataset_generator_fun_train, num_parallel_calls=tf.data.AUTOTUNE)\
    .shuffle(BUFFER_SIZE)\
    .batch(BATCH_SIZE, drop_remainder=True)\
    .map(map_fun_with_time_batchwise, num_parallel_calls=tf.data.AUTOTUNE)\
    .cache()\
    .prefetch(tf.data.AUTOTUNE)

dataset_test_optimized = tf.data.Dataset.range(1).interleave(dataset_generator_fun_test, num_parallel_calls=tf.data.AUTOTUNE)\
    .batch(BATCH_SIZE, drop_remainder=True)\
    .map(map_fun_test_with_time_batchwise, num_parallel_calls=tf.data.AUTOTUNE)\
    .cache()\
    .prefetch(tf.data.AUTOTUNE)
```

@map\_decorator

def map\_fun\_with\_time\_batchwise(steps, times, values, image, label):

*# sleep to avoid concurrency issue*

time.sleep(0.05)

map\_enter = time.perf\_counter()

image = tf.reshape(image, [tf.shape(image)[0], IMAGE\_DEPTH, IMAGE\_HEIGHT, IMAGE\_WIDTH])

image = tf.divide(tf.cast(tf.transpose(image, [0, 2, 3, 1]), tf.float32), 255.0)

label = tf.one\_hot(label, 10)

distorted\_image = tf.image.random\_crop(image, [tf.shape(image)[0], IMAGE\_SIZE\_CROPPED, IMAGE\_SIZE\_CROPPED, IMAGE\_DEPTH])

distorted\_image = tf.image.random\_flip\_left\_right(distorted\_image)

distorted\_image = tf.image.random\_brightness(distorted\_image, max\_delta=63)

distorted\_image = tf.image.random\_contrast(distorted\_image, lower=0.2, upper=1.8)

distorted\_image = tf.image.per\_image\_standardization(distorted\_image)

# Assignment

- Goal
  - Try some the input transformation mentioned above (e.g. shuffle, batch, repeat, `map(random_crop, random_flip_left_right, ...)`) but without optimization terms (e.g. `prefetch`, `cache`, `num_parallel_calls`), comparing the performance to the no input transformation
  - Retrain your model with optimized terms, compare the time consuming
  - Training both models above for at least 3 epochs
  - Briefly summarize what you did and explain the performance results (accuracy and time consuming)
- Deadline: 2025/10/23 (Thr) 00:00